

# capstone-credit card fraud detection

May 2, 2019

## 1 Introduction

The datasets contains transactions made by credit cards in September 2013 by european cardholders. This dataset presents transactions that occurred in two days, where we have 492 frauds out of 284,807 transactions. The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions.

It contains only numerical input variables which are the result of a PCA transformation. Unfortunately, due to confidentiality issues, we cannot provide the original features and more background information about the data. Features V1, V2, ... V28 are the principal components obtained with PCA, the only features which have not been transformed with PCA are 'Time' and 'Amount'. Feature 'Time' contains the seconds elapsed between each transaction and the first transaction in the dataset. The feature 'Amount' is the transaction Amount, this feature can be used for example-dependant cost-sensitive learning. Feature 'Class' is the response variable and it takes value 1 in case of fraud and 0 otherwise.

In [1]: *## Using IBM Cloud*

```
# import types
# import pandas as pd
# from boto3.client import Config
# import ibm_boto3

# def __iter__(self): return 0

## @hidden_cell
## The following code accesses a file in your IBM Cloud Object Storage. It includes y
## You might want to remove those credentials before you share your notebook.
# client_dcc662e2032147378976f288b5e93fcc = ibm_boto3.client(service_name='s3',
#     ibm_api_key_id='pGKI3UAgcg9d2BDY9-eXqh5g00qv-R-AJryFg9j6-0gJ',
#     ibm_auth_endpoint="https://iam.bluemix.net/oidc/token",
#     config=Config(signature_version='oauth'),
#     endpoint_url='https://s3-api.us-geo.objectstorage.service.networklayer.com')

# body = client_dcc662e2032147378976f288b5e93fcc.get_object(Bucket='default-donotdelet
## add missing __iter__ method, so pandas accepts body as file-like object
# if not hasattr(body, "__iter__"): body.__iter__ = types.MethodType(__iter__, body )
```

```
# df = pd.read_csv(body)
# df.head()
```

```
Out[1]:
```

|   | Time | V1        | V2        | V3       | V4        | V5        | V6        | V7        | \ |
|---|------|-----------|-----------|----------|-----------|-----------|-----------|-----------|---|
| 0 | 0.0  | -1.359807 | -0.072781 | 2.536347 | 1.378155  | -0.338321 | 0.462388  | 0.239599  |   |
| 1 | 0.0  | 1.191857  | 0.266151  | 0.166480 | 0.448154  | 0.060018  | -0.082361 | -0.078803 |   |
| 2 | 1.0  | -1.358354 | -1.340163 | 1.773209 | 0.379780  | -0.503198 | 1.800499  | 0.791461  |   |
| 3 | 1.0  | -0.966272 | -0.185226 | 1.792993 | -0.863291 | -0.010309 | 1.247203  | 0.237609  |   |
| 4 | 2.0  | -1.158233 | 0.877737  | 1.548718 | 0.403034  | -0.407193 | 0.095921  | 0.592941  |   |

|   | V8        | V9        | ... | V21       | V22       | V23       | V24       | \ |
|---|-----------|-----------|-----|-----------|-----------|-----------|-----------|---|
| 0 | 0.098698  | 0.363787  | ... | -0.018307 | 0.277838  | -0.110474 | 0.066928  |   |
| 1 | 0.085102  | -0.255425 | ... | -0.225775 | -0.638672 | 0.101288  | -0.339846 |   |
| 2 | 0.247676  | -1.514654 | ... | 0.247998  | 0.771679  | 0.909412  | -0.689281 |   |
| 3 | 0.377436  | -1.387024 | ... | -0.108300 | 0.005274  | -0.190321 | -1.175575 |   |
| 4 | -0.270533 | 0.817739  | ... | -0.009431 | 0.798278  | -0.137458 | 0.141267  |   |

|   | V25       | V26       | V27       | V28       | Amount | Class |
|---|-----------|-----------|-----------|-----------|--------|-------|
| 0 | 0.128539  | -0.189115 | 0.133558  | -0.021053 | 149.62 | 0     |
| 1 | 0.167170  | 0.125895  | -0.008983 | 0.014724  | 2.69   | 0     |
| 2 | -0.327642 | -0.139097 | -0.055353 | -0.059752 | 378.66 | 0     |
| 3 | 0.647376  | -0.221929 | 0.062723  | 0.061458  | 123.50 | 0     |
| 4 | -0.206010 | 0.502292  | 0.219422  | 0.215153  | 69.99  | 0     |

[5 rows x 31 columns]

```
In [1]: # running on local machine
import pandas as pd
df = pd.read_csv('creditcardfraud/creditcard.csv')

In [2]: # it has 28 features plus time and amount variables
df.columns
```

```
Out[2]: Index(['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10',
              'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20',
              'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount',
              'Class'],
              dtype='object')
```

```
In [3]: # take a look at the statistics of these features
df.describe()
```

```
Out[3]:
```

|       | Time          | V1            | V2            | V3            | V4            | \ |
|-------|---------------|---------------|---------------|---------------|---------------|---|
| count | 284807.000000 | 2.848070e+05  | 2.848070e+05  | 2.848070e+05  | 2.848070e+05  |   |
| mean  | 94813.859575  | 3.919560e-15  | 5.688174e-16  | -8.769071e-15 | 2.782312e-15  |   |
| std   | 47488.145955  | 1.958696e+00  | 1.651309e+00  | 1.516255e+00  | 1.415869e+00  |   |
| min   | 0.000000      | -5.640751e+01 | -7.271573e+01 | -4.832559e+01 | -5.683171e+00 |   |
| 25%   | 54201.500000  | -9.203734e-01 | -5.985499e-01 | -8.903648e-01 | -8.486401e-01 |   |
| 50%   | 84692.000000  | 1.810880e-02  | 6.548556e-02  | 1.798463e-01  | -1.984653e-02 |   |

|     |               |              |              |              |              |
|-----|---------------|--------------|--------------|--------------|--------------|
| 75% | 139320.500000 | 1.315642e+00 | 8.037239e-01 | 1.027196e+00 | 7.433413e-01 |
| max | 172792.000000 | 2.454930e+00 | 2.205773e+01 | 9.382558e+00 | 1.687534e+01 |

|       |               |               |               |               |               |
|-------|---------------|---------------|---------------|---------------|---------------|
|       | V5            | V6            | V7            | V8            | V9 \          |
| count | 2.848070e+05  | 2.848070e+05  | 2.848070e+05  | 2.848070e+05  | 2.848070e+05  |
| mean  | -1.552563e-15 | 2.010663e-15  | -1.694249e-15 | -1.927028e-16 | -3.137024e-15 |
| std   | 1.380247e+00  | 1.332271e+00  | 1.237094e+00  | 1.194353e+00  | 1.098632e+00  |
| min   | -1.137433e+02 | -2.616051e+01 | -4.355724e+01 | -7.321672e+01 | -1.343407e+01 |
| 25%   | -6.915971e-01 | -7.682956e-01 | -5.540759e-01 | -2.086297e-01 | -6.430976e-01 |
| 50%   | -5.433583e-02 | -2.741871e-01 | 4.010308e-02  | 2.235804e-02  | -5.142873e-02 |
| 75%   | 6.119264e-01  | 3.985649e-01  | 5.704361e-01  | 3.273459e-01  | 5.971390e-01  |
| max   | 3.480167e+01  | 7.330163e+01  | 1.205895e+02  | 2.000721e+01  | 1.559499e+01  |

|       |     |               |               |               |               |
|-------|-----|---------------|---------------|---------------|---------------|
|       | ... | V21           | V22           | V23           | V24 \         |
| count | ... | 2.848070e+05  | 2.848070e+05  | 2.848070e+05  | 2.848070e+05  |
| mean  | ... | 1.537294e-16  | 7.959909e-16  | 5.367590e-16  | 4.458112e-15  |
| std   | ... | 7.345240e-01  | 7.257016e-01  | 6.244603e-01  | 6.056471e-01  |
| min   | ... | -3.483038e+01 | -1.093314e+01 | -4.480774e+01 | -2.836627e+00 |
| 25%   | ... | -2.283949e-01 | -5.423504e-01 | -1.618463e-01 | -3.545861e-01 |
| 50%   | ... | -2.945017e-02 | 6.781943e-03  | -1.119293e-02 | 4.097606e-02  |
| 75%   | ... | 1.863772e-01  | 5.285536e-01  | 1.476421e-01  | 4.395266e-01  |
| max   | ... | 2.720284e+01  | 1.050309e+01  | 2.252841e+01  | 4.584549e+00  |

|       |               |               |               |               |               |
|-------|---------------|---------------|---------------|---------------|---------------|
|       | V25           | V26           | V27           | V28           | Amount \      |
| count | 2.848070e+05  | 2.848070e+05  | 2.848070e+05  | 2.848070e+05  | 284807.000000 |
| mean  | 1.453003e-15  | 1.699104e-15  | -3.660161e-16 | -1.206049e-16 | 88.349619     |
| std   | 5.212781e-01  | 4.822270e-01  | 4.036325e-01  | 3.300833e-01  | 250.120109    |
| min   | -1.029540e+01 | -2.604551e+00 | -2.256568e+01 | -1.543008e+01 | 0.000000      |
| 25%   | -3.171451e-01 | -3.269839e-01 | -7.083953e-02 | -5.295979e-02 | 5.600000      |
| 50%   | 1.659350e-02  | -5.213911e-02 | 1.342146e-03  | 1.124383e-02  | 22.000000     |
| 75%   | 3.507156e-01  | 2.409522e-01  | 9.104512e-02  | 7.827995e-02  | 77.165000     |
| max   | 7.519589e+00  | 3.517346e+00  | 3.161220e+01  | 3.384781e+01  | 25691.160000  |

|       |               |
|-------|---------------|
|       | Class         |
| count | 284807.000000 |
| mean  | 0.001727      |
| std   | 0.041527      |
| min   | 0.000000      |
| 25%   | 0.000000      |
| 50%   | 0.000000      |
| 75%   | 0.000000      |
| max   | 1.000000      |

[8 rows x 31 columns]

```
In [4]: # see if there is missing values
df.isnull().sum()
```

```
Out[4]: Time      0
```

```

V1      0
V2      0
V3      0
V4      0
V5      0
V6      0
V7      0
V8      0
V9      0
V10     0
V11     0
V12     0
V13     0
V14     0
V15     0
V16     0
V17     0
V18     0
V19     0
V20     0
V21     0
V22     0
V23     0
V24     0
V25     0
V26     0
V27     0
V28     0
Amount  0
Class   0
dtype: int64

```

```

In [5]: tmp = df.Class.value_counts()
        print('No frauds: ', tmp[0])
        print('Frauds: ', tmp[1])

```

```

No frauds: 284315
Frauds: 492

```

```

In [6]: import seaborn as sns
        import matplotlib.pyplot as plt
        import warnings
        warnings.filterwarnings("ignore")
        %matplotlib inline


```

```

In [7]: # visualize the classes distribution
        sns.countplot('Class', data=df, palette="muted")

```

Out[7]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1a13f5cd3c8>



As we can see from the above plot, our original dataset is extremely imbalanced. Most of the transactions are non-fraud. If we use this dataset without any modification to train our classifiers, then we will most certainly overfit.


```
In [8]: # lets also examine the distributions of time and amount features
fig, ax = plt.subplots(1, 2, figsize=(18,4))

amount_val = df['Amount'].values
time_val = df['Time'].values

sns.distplot(amount_val, ax=ax[0], color='r')
ax[0].set_title('Distribution of Transaction Amount', fontsize=14)
ax[0].set_xlim([min(amount_val), max(amount_val)])

sns.distplot(time_val, ax=ax[1], color='b')
ax[1].set_title('Distribution of Transaction Time', fontsize=14)
ax[1].set_xlim([min(time_val), max(time_val)])
```

Out[8]: (0.0, 172792.0)



As shown in the above plots, both Time and Transaction Amount are needed to be scaled in order to improve model performances.

## 1.1 Scaling Time and Amount

```
In [9]: from sklearn.preprocessing import RobustScaler
        scaler = RobustScaler() # robust scaler is less prone to outliers
        df['scaled_time'] = scaler.fit_transform(df.Time.values.reshape(-1,1))
        df['scaled_amount'] = scaler.fit_transform(df.Amount.values.reshape(-1,1))
        df.drop(['Time', 'Amount'], axis=1, inplace=True)
```

```
In [10]: df.head()
```

```
Out[10]:
```

|   | V1        | V2        | V3       | V4        | V5        | V6        | V7        | \ |
|---|-----------|-----------|----------|-----------|-----------|-----------|-----------|---|
| 0 | -1.359807 | -0.072781 | 2.536347 | 1.378155  | -0.338321 | 0.462388  | 0.239599  |   |
| 1 | 1.191857  | 0.266151  | 0.166480 | 0.448154  | 0.060018  | -0.082361 | -0.078803 |   |
| 2 | -1.358354 | -1.340163 | 1.773209 | 0.379780  | -0.503198 | 1.800499  | 0.791461  |   |
| 3 | -0.966272 | -0.185226 | 1.792993 | -0.863291 | -0.010309 | 1.247203  | 0.237609  |   |
| 4 | -1.158233 | 0.877737  | 1.548718 | 0.403034  | -0.407193 | 0.095921  | 0.592941  |   |

|   | V8        | V9        | V10       | ... | V22       | V23       | V24       | V25       | \ |
|---|-----------|-----------|-----------|-----|-----------|-----------|-----------|-----------|---|
| 0 | 0.098698  | 0.363787  | 0.090794  | ... | 0.277838  | -0.110474 | 0.066928  | 0.128539  |   |
| 1 | 0.085102  | -0.255425 | -0.166974 | ... | -0.638672 | 0.101288  | -0.339846 | 0.167170  |   |
| 2 | 0.247676  | -1.514654 | 0.207643  | ... | 0.771679  | 0.909412  | -0.689281 | -0.327642 |   |
| 3 | 0.377436  | -1.387024 | -0.054952 | ... | 0.005274  | -0.190321 | -1.175575 | 0.647376  |   |
| 4 | -0.270533 | 0.817739  | 0.753074  | ... | 0.798278  | -0.137458 | 0.141267  | -0.206010 |   |


|   | V26       | V27       | V28       | Class | scaled_time | scaled_amount |
|---|-----------|-----------|-----------|-------|-------------|---------------|
| 0 | -0.189115 | 0.133558  | -0.021053 | 0     | -0.994983   | 1.783274      |
| 1 | 0.125895  | -0.008983 | 0.014724  | 0     | -0.994983   | -0.269825     |
| 2 | -0.139097 | -0.055353 | -0.059752 | 0     | -0.994972   | 4.983721      |
| 3 | -0.221929 | 0.062723  | 0.061458  | 0     | -0.994972   | 1.418291      |
| 4 | 0.502292  | 0.219422  | 0.215153  | 0     | -0.994960   | 0.670579      |

```
[5 rows x 31 columns]
```

```
In [11]: # lets take a look at the correlation plot to examine features relation to classes
        plt.figure(figsize=(15,10))
```

```
plt.title("Imbalanced Correlation Matrix")
sns.heatmap(df.corr(), cmap='coolwarm_r', annot_kws={'size':20})
```

Out[11]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1a13f65d630>



## 1.2 subsample data with equal observation from both classes and explore correlation matrix again

```
In [12]: # shuffle the data
df = df.sample(frac=1)

frauds = df.loc[df.Class==1]
non_frauds = df.loc[df.Class==0][:429]
new_df = pd.concat([frauds, non_frauds])
new_df = new_df.sample(frac=1, random_state=2019)
new_df.head()
```

Out[12]:

|        | V1        | V2        | V3        | V4        | V5        | V6        | V7        | \ |
|--------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|---|
| 140394 | -6.961676 | 3.882633  | -2.892660 | -0.057403 | -2.866139 | -0.721680 | -2.084298 |   |
| 6472   | 1.023874  | 2.001485  | -4.769752 | 3.819195  | -1.271754 | -1.734662 | -3.059245 |   |
| 259002 | 2.029039  | -0.928822 | 0.048696  | -0.724177 | -1.359196 | -0.319260 | -1.304253 |   |

```

77099 -0.075483  1.812355 -2.566981  4.127549 -1.628532 -0.805895 -3.390135
105980  1.524484 -1.053841  0.497467 -1.442412 -1.524713 -0.646696 -1.110334

```

```

          V8          V9          V10  ...          V22          V23          V24  \
140394  3.463034  0.648002  1.615223  ... -0.418317  0.422675  0.056979
6472    0.889805  0.415382 -3.955812  ... -0.054196  0.709654 -0.372216
259002  0.183088  2.312593 -0.422458  ...  0.667772  0.218061 -0.387726
77099   1.019353 -2.451251 -3.555835  ...  0.270471 -0.143624  0.013566
105980 -0.120683 -1.635291  1.425733  ... -0.843161  0.117493 -0.185511

```

```

          V25          V26          V27          V28  Class  scaled_time  \
140394  0.820109  0.273678  0.204325 -0.354226      0    -0.011666
6472    -2.032068  0.366778  0.395171  0.020206      1    -0.904052
259002 -0.563044  0.713271 -0.003171 -0.047048      0     0.871967
77099   0.634203  0.213693  0.773625  0.387434      1    -0.326660
105980  0.179495 -0.377980  0.038188  0.024832      0    -0.175319

```

```

          scaled_amount
140394          0.156082
6472         -0.293440
259002         -0.293440
77099         -0.237546
105980         -0.033536

```

```
[5 rows x 31 columns]
```


```

In [13]: plt.figure(figsize=(15,10))
plt.title("Balanced Correlation Matrix")
sns.heatmap(new_df.corr(), cmap='coolwarm_r', annot_kws={'size':20})

```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x1a141c6fac8>
```





By comparing the two correlation matrix heatmaps, we can see that there is a noticeable difference in correlations using balanced and imbalanced datasets, and we can observe more correlated relations when using balanced datasets. Therefore, it suggests that we might want to use a more balanced dataset to train a robust classifier.

## 2 TSNE Algorithm to cluster our new balanced sample

```
In [14]: from sklearn.manifold import TSNE
import matplotlib.patches as mpatches
from mpl_toolkits.mplot3d import Axes3D

def plot_tsne(dim=2):
    # New_df is from the random undersample data (fewer instances)
    X = new_df.drop('Class', axis=1)
    y = new_df['Class']
    blue_patch = mpatches.Patch(color='#0A0AFF', label='No Fraud')
    red_patch = mpatches.Patch(color='#AF0000', label='Fraud')

    # T-SNE Implementation
    X_reduced_tsne = TSNE(n_components=dim, random_state=2019).fit_transform(X.values)
    if dim == 3:
```

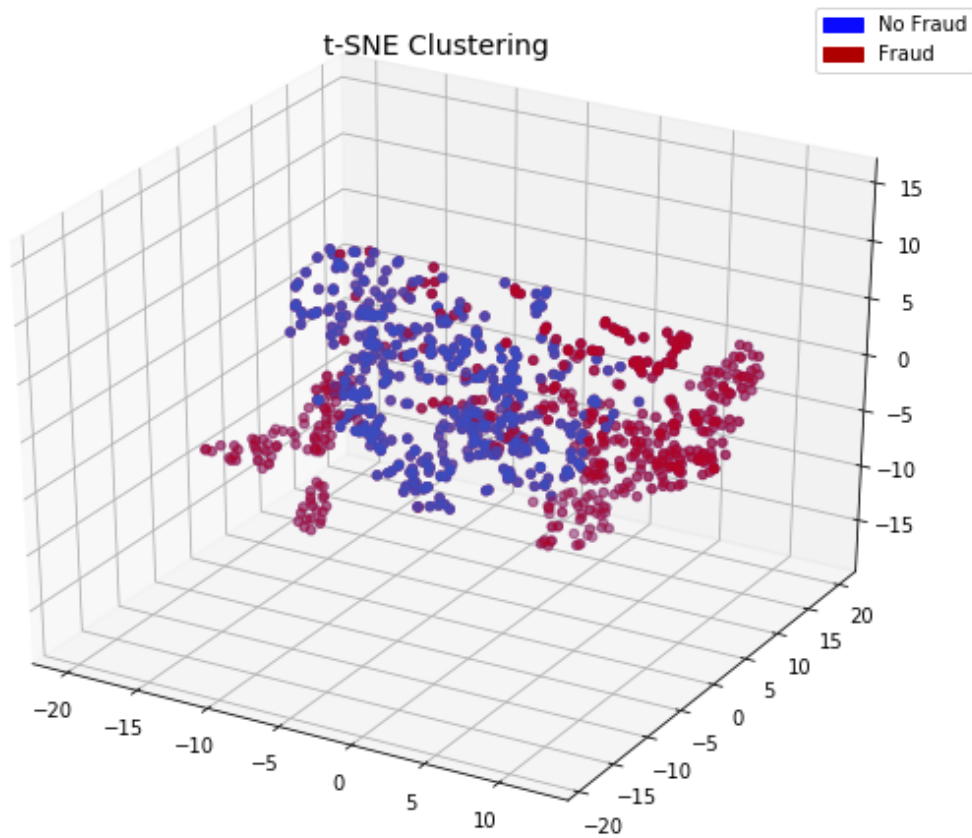
```

fig = plt.figure(figsize=(10, 8))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(X_reduced_tsne[:,0], X_reduced_tsne[:,1], X_reduced_tsne[:,2], c=(
ax.scatter(X_reduced_tsne[:,0], X_reduced_tsne[:,1], X_reduced_tsne[:,2], c=(
ax.set_title('t-SNE Clustering', fontsize=14)
ax.grid(True)
ax.legend(handles=[blue_patch, red_patch])
elif dim == 2:
plt.figure(figsize=(12,8))
plt.scatter(X_reduced_tsne[:,0], X_reduced_tsne[:,1], c=(y == 0), cmap='coolw
plt.scatter(X_reduced_tsne[:,0], X_reduced_tsne[:,1], c=(y == 1), cmap='coolw
plt.title('t-SNE Clustering', fontsize=14)
plt.grid(True)
plt.legend(handles=[blue_patch, red_patch])
else:
    return print('invalid input dimension')

```

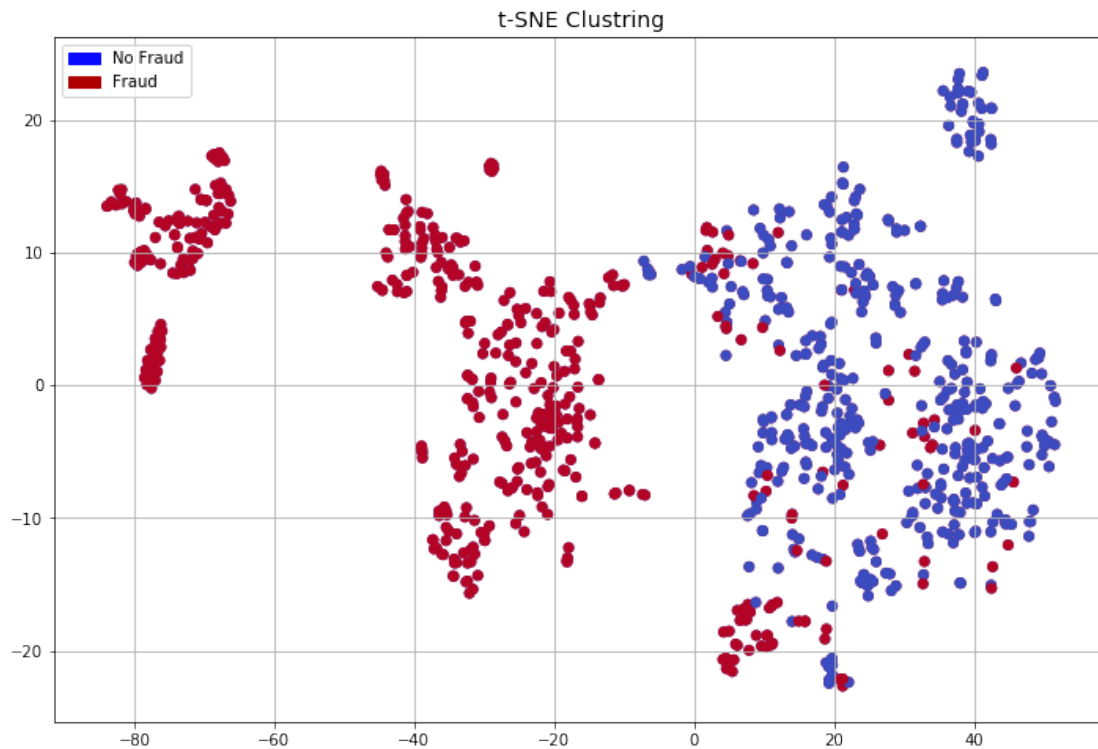
In [15]: %time plot\_tsne(3)

Wall time: 16.2 s



```
In [16]: %time plot_tsne(2)
```

Wall time: 5.23 s



### 3 SMOTE (Synthetic Minority Over-sampling Technique)

We will use SMOTE to address to imbalanced datasets

- SMOTE will generate synthetic points from the minority class in order to make our training set more balance
- SMOTE will picks the distance between the closest neighbors of the minority class, in between these distances it creates synthetic points
- More information is obtained

```
In [17]: !pip install imbalanced-learn
```

Collecting imbalanced-learn

Downloading <https://files.pythonhosted.org/packages/e5/4c/7557e1c2e791bd43878f8c82065bddc5798>  
Requirement already satisfied: scipy>=0.13.3 in c:\users\michael\anaconda3\lib\site-packages (0.13.3)  
Requirement already satisfied: scikit-learn>=0.20 in c:\users\michael\anaconda3\lib\site-packages (0.20)  
Requirement already satisfied: numpy>=1.8.2 in c:\users\michael\anaconda3\lib\site-packages (1.8.2)  
Installing collected packages: imbalanced-learn  
Successfully installed imbalanced-learn-0.4.3

## 4 Train and Evaluate

```
In [18]: def plot_cm(cm):
        cm = pd.DataFrame(cm)
        sns.heatmap(cm, annot=True, fmt='.1f', cmap='YlGnBu', linewidths=5, vmax=500)
        plt.yticks(rotation=0)
        plt.show()

In [19]: from sklearn.model_selection import StratifiedShuffleSplit
        from imblearn.over_sampling import SMOTE
        from sklearn.metrics import precision_score, recall_score, f1_score, roc_auc_score, a
        from sklearn.model_selection import GridSearchCV
        import time

        # split original datasets into features and label data
        X = df.drop(['Class'], axis=1).values
        y = df.Class.values

        def cross_val(clf, grid_param=None, k=5, test=0.5, cv_test=0.2, plot=True):

            sss = StratifiedShuffleSplit(n_splits=1, test_size=test, random_state=2019)
            for trn_idx, tes_idx in sss.split(X,y):
                time1 = time.time()
                x_trn, x_tes = X[trn_idx], X[tes_idx]
                y_trn, y_tes = y[trn_idx], y[tes_idx]
                x_res, y_res = SMOTE(sampling_strategy='minority').fit_resample(x_trn, y_trn)
                auc = []

                k = 1 if grid_param else k # if use GridSearchCV, use one fold becuae GridSe

            for fold, (tr_idx, te_idx) in enumerate(StratifiedShuffleSplit(n_splits=k, tes
                Xtrain, Xtest = x_res[tr_idx], x_res[te_idx]
                ytrain, ytest = y_res[tr_idx], y_res[te_idx]
                if grid_param:
                    cv = GridSearchCV(clf, grid_param, scoring='roc_auc', verbose=1, n_jo
                    cv.fit(Xtrain, ytrain)
                    score = roc_auc_score(ytest, cv.predict(Xtest))
                    print(fold, ': ', score)
                    auc.append(score)
                else:
                    clf.fit(Xtrain, ytrain)
                    score = roc_auc_score(ytest, clf.predict(Xtest))
                    print(fold, ': ', score)
                    auc.append(score)
            print(k, 'cv score: ', sum(auc)/len(auc))

            if grid_param:
```

```

        preds = cv.predict(x_tes)
        print('test score: ', roc_auc_score(y_tes, preds))
        print('best parameters: ', cv.best_params_)
    else:
        preds = clf.predict(x_tes)
        print('test score: ', roc_auc_score(y_tes, preds))

    print('Completion time:', time.time()-time1, 'seconds')

    if plot:
        cm = confusion_matrix(y_tes, preds)
        plot_cm(cm)
        print(classification_report(y_tes, preds))

    return cv if grid_param else clf

```

## 5 Model Selection

We will define a number of classifiers and train them on just one fold training dataset, and we will select just a small subset of models for tuning based on the ROC\_AUC\_SCORE and time efficiency.

```

In [20]: from sklearn.linear_model import LogisticRegression
        from sklearn.neighbors import KNeighborsClassifier
        from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
        from sklearn.svm import SVC

```

```

classifiers = {
    "LogisticRegression": LogisticRegression(),
    "RandomForestClassifier": RandomForestClassifier(),
    "kNieghborsClassifier": KNeighborsClassifier(),
    "SVC": SVC(),
    "GradientBoostingClassifier": GradientBoostingClassifier()
}

```

```


clfs = {}
for clf in classifiers:
    print(clf)
    clfs.setdefault(clf, cross_val(classifiers[clf], k=3))

```

```

LogisticRegression
0 : 0.9573004919877597
1 : 0.9576002290353881
2 : 0.9565618770830947
3 cv score:  0.9571541993687475
test score:  0.9223681551169962
Completion time: 13.247580289840698 seconds

```




|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 0.98   | 0.99     | 142158  |
| 1            | 0.07      | 0.87   | 0.12     | 246     |
| micro avg    | 0.98      | 0.98   | 0.98     | 142404  |
| macro avg    | 0.53      | 0.92   | 0.56     | 142404  |
| weighted avg | 1.00      | 0.98   | 0.99     | 142404  |

```

RandomForestClassifier
0 : 0.9998768992684299
1 : 0.9999120678133024
2 : 0.9999472425436129
3 cv score: 0.999912069875115
test score: 0.8881340034225058
Completion time: 31.9435932636261 seconds

```




|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 142158  |
| 1            | 0.90      | 0.78   | 0.83     | 246     |
| micro avg    | 1.00      | 1.00   | 1.00     | 142404  |
| macro avg    | 0.95      | 0.89   | 0.92     | 142404  |
| weighted avg | 1.00      | 1.00   | 1.00     | 142404  |

```

kNeighborsClassifier
0 : 0.9991207090602138
1 : 0.9990327459463262
2 : 0.9990855374226224
3 cv score: 0.999079664143054
test score: 0.9238893069511457
Completion time: 520.0365943908691 seconds

```



|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 142158  |
| 1            | 0.45      | 0.85   | 0.59     | 246     |
| micro avg    | 1.00      | 1.00   | 1.00     | 142404  |
| macro avg    | 0.72      | 0.92   | 0.79     | 142404  |
| weighted avg | 1.00      | 1.00   | 1.00     | 142404  |

SVC

0 : 0.9994196679797411

1 : 0.9992437831944005


2 : 0.9992789814293754

3 cv score: 0.9993141442011724

test score: 0.8285824361008142

Completion time: 768.6259486675262 seconds





|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 142158  |
| 1            | 0.45      | 0.66   | 0.54     | 246     |
| micro avg    | 1.00      | 1.00   | 1.00     | 142404  |
| macro avg    | 0.73      | 0.83   | 0.77     | 142404  |
| weighted avg | 1.00      | 1.00   | 1.00     | 142404  |

GradientBoostingClassifier

0 : 0.9926841466006605


1 : 0.9926490188796788

2 : 0.992561034735302

3 cv score: 0.9926314000718804

test score: 0.9257490548990662

Completion time: 222.79033279418945 seconds



|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 0.99   | 1.00     | 142158  |
| 1            | 0.19      | 0.86   | 0.31     | 246     |
| micro avg    | 0.99      | 0.99   | 0.99     | 142404  |
| macro avg    | 0.60      | 0.93   | 0.66     | 142404  |
| weighted avg | 1.00      | 0.99   | 1.00     | 142404  |

We should look at our trained models, and determine how to set values for parameters tuning

```
In [21]: import pickle
         for clf in clfs:
             pickle.dump(clfs[clf], open('models/'+clf+'_intial.sav', 'wb'))
```

```
In [29]: !dir models
```

```
Volume in drive C has no label.
Volume Serial Number is 244D-FC98
```

```
Directory of C:\Users\michael\Desktop\coursera\Advanced Data Science with IBM Specialization\
```

```
05/02/2019  05:17 PM    <DIR>          .
05/02/2019  05:17 PM    <DIR>          ..
```

```

05/02/2019  04:16 PM          135,848 GradientBoostingClassifier_intial.sav
05/02/2019  05:17 PM          243,511 GradientBoostingClassifier_tuned.sav
05/02/2019  04:16 PM      61,512,819 kNeighborsClassifier_intial.sav
05/02/2019  04:16 PM           1,042 LogisticRegression_intial.sav
05/02/2019  05:17 PM           5,617 LogisticRegression_tuned.sav
05/02/2019  04:16 PM      721,425 RandomForestClassifier_intial.sav
05/02/2019  05:17 PM    1,496,013 RandomForestClassifier_tuned.sav
05/02/2019  04:16 PM    1,164,080 SVC_intial.sav
           8 File(s)      65,280,355 bytes
           2 Dir(s)  865,668,763,648 bytes free

```

From the above results, we can see that GradientBoostingClassifier, RandomForestClassifier, and LogisticClassifier did yield a relatively good results with a very fast completeion time, therefore, we can further fine tune these 3 models, and compare how they perform.

## 6 Model Tuning

In [23]: *# define parameters values*

```

log_params = {"penalty": ['l1', 'l2'], 'C': [0.001, 0.01, 0.1, 1, 10, 100], "n_jobs":
rf_params = {"min_samples_leaf": [1,3,5,10,25], "max_features": [1, 0.5, 'log2', 'sqrt'],
gb_params = {"min_samples_leaf": [100, 500], "max_depth": [5, 8], "subsample": [0.3, 0.5, 0.7, 0.9]}

params = {
    "LogisticRegression": log_params,
    "RandomForestClassifier": rf_params,
    "GradientBoostingClassifier": gb_params
}

classifiers = {
    "LogisticRegression": LogisticRegression(),
    "RandomForestClassifier": RandomForestClassifier(),
    "GradientBoostingClassifier": GradientBoostingClassifier()
}

```

```

In [25]: clfs_tuned = {}
        for clf in classifiers:
            print('Tuning', clf)
            clfs_tuned.setdefault(clf, cross_val(classifiers[clf], grid_param=params[clf]))

```

Tuning LogisticRegression

Fitting 2 folds for each of 12 candidates, totalling 24 fits

```


[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 24 out of 24 | elapsed: 19.2s finished

```

```

0 : 0.9559990956246349
1 cv score: 0.9559990956246349
test score: 0.924386606589233
best parameters: {'C': 100, 'n_jobs': -1, 'penalty': 'l2'}
Completion time: 25.183666944503784 seconds

```



|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 0.98   | 0.99     | 142158  |
| 1            | 0.07      | 0.87   | 0.12     | 246     |
| micro avg    | 0.98      | 0.98   | 0.98     | 142404  |
| macro avg    | 0.53      | 0.92   | 0.56     | 142404  |
| weighted avg | 1.00      | 0.98   | 0.99     | 142404  |

```

Tuning RandomForestClassifier
Fitting 2 folds for each of 20 candidates, totalling 40 fits

```

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 40 out of 40 | elapsed: 51.6s finished


```

```

0 : 0.9997537972997723
1 cv score: 0.9997537972997723

```

test score: 0.8819063055569567  
 best parameters: {'max\_features': 1, 'min\_samples\_leaf': 5, 'n\_jobs': -1}  
 Completion time: 53.95374536514282 seconds




|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 142158  |
| 1            | 0.76      | 0.76   | 0.76     | 246     |
| micro avg    | 1.00      | 1.00   | 1.00     | 142404  |
| macro avg    | 0.88      | 0.88   | 0.88     | 142404  |
| weighted avg | 1.00      | 1.00   | 1.00     | 142404  |

Tuning GradientBoostingClassifier  
 Fitting 2 folds for each of 48 candidates, totalling 96 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.  
 [Parallel(n\_jobs=-1)]: Done 34 tasks | elapsed: 12.7min  
 [Parallel(n\_jobs=-1)]: Done 96 out of 96 | elapsed: 36.5min finished

0 : 0.9997362127180642  
 1 cv score: 0.9997362127180642  
 test score: 0.9081743124019683

best parameters: {'learning\_rate': 0.5, 'max\_depth': 5, 'min\_samples\_leaf': 500, 'n\_estimators': 1000}  
 Completion time: 2246.9703459739685 seconds



|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 142158  |
| 1            | 0.66      | 0.82   | 0.73     | 246     |
| micro avg    | 1.00      | 1.00   | 1.00     | 142404  |
| macro avg    | 0.83      | 0.91   | 0.87     | 142404  |
| weighted avg | 1.00      | 1.00   | 1.00     | 142404  |

As a result of tuning our models' parameters, we can see that all of our models have a small improvement in terms of ROC\_AUC score, and the LogisticClassifier has the best performance, therefore, we can decide to use LogisticClassifier as our final model for deployment.

```
In [27]: # saving models
         for clf in clfs_tuned:
             pickle.dump(clfs_tuned[clf], open('models/'+clf+'_tuned.sav', 'wb'))
```

```
In [28]: !dir models
```

```
Volume in drive C has no label.
Volume Serial Number is 244D-FC98
```

Directory of C:\Users\michael\Desktop\coursera\Advanced Data Science with IBM Specialization\

```
05/02/2019 05:17 PM <DIR> .
05/02/2019 05:17 PM <DIR> ..
05/02/2019 04:16 PM      135,848 GradientBoostingClassifier_intial.sav
05/02/2019 05:17 PM      243,511 GradientBoostingClassifier_tuned.sav
05/02/2019 04:16 PM    61,512,819 kNieghborsClassifier_intial.sav
05/02/2019 04:16 PM      1,042 LogisticRegression_intial.sav
05/02/2019 05:17 PM      5,617 LogisticRegression_tuned.sav
05/02/2019 04:16 PM      721,425 RandomForestClassifier_intial.sav
05/02/2019 05:17 PM    1,496,013 RandomForestClassifier_tuned.sav
05/02/2019 04:16 PM    1,164,080 SVC_intial.sav
      8 File(s)      65,280,355 bytes
      2 Dir(s)  865,669,206,016 bytes free
```